

Paul J. Yoon

+1 (713) 320-1032 | pauljy@stanford.edu | [linkedin.com/in/pauljinyoon](https://www.linkedin.com/in/pauljinyoon)

EDUCATION

STANFORD UNIVERSITY

Bachelor of Science, Mathematics. Minor in Music

Expected Graduation: Jun 2027

GPA: 3.8/4.0

Relevant Coursework: Linear Algebra, Differential, and Integral Calculus of Several Variables, Computer Organization and Systems, Real Analysis, Probability Theory for Computer Scientists, Machine Learning, Math for Machine Learning

PROFESSIONAL EXPERIENCE

STANFORD DEPARTMENT OF MATHEMATICS

Student Researcher

Stanford, CA

Jun 2025– Aug 2025

- Blueprinted the full proof of Wigner’s Semicircle law into a 20+ page LaTeX blueprint bridging probability theory, combinatorics, and measure theory
- Implemented and machine-verified dozens of core lemmas and constructions in Lean, preparing contributions to Lean’s mathlib library
- Collaborated with a three person research team using Git version control workflows to maintain a shared codebase

SUNDIAL

Data Science Intern

Palo Alto, CA

Jul 2024 – Sep 2024

- Sundial is a Series A startup building an automated data science and product analytics tool founded by two early Meta executives, one of whom went to Sequoia Capital before founding Sundial
- Developed classifier identifying fraudulent users using behavior-based thresholds, reducing false positives compared to existing process by 50% and boosting overall detection accuracy by 75%
- Created, trained, and tuned a time series seasonality model, outperforming existing model by 120% as measured by mean absolute percentage error (MAPE)

PROJECTS

Timestamping Video Game Eliminations with Computer Vision

Python, LaTeX

Apr 2025 – Jun 2025

- Curated a custom detection dataset (165 events, ~ 1,650 frames) with color-jitter augmentation strategies
- Fine-tuned YOLOv8-nano (2.5M params) on <200 images, achieving 0.61 F1 and 0.49s mean temporal error, doubling precision compared to a baseline template matcher
- Engineered a lightweight inference pipeline (OpenCV + ffmpeg) that processes a 14-minute VOD in 1.5 minutes on CPU, generating highlight clips with sub-frame accuracy

An Exploratory Analysis of Feature Representation in Music Source Separation

Python, LaTeX

Jan 2025 – Mar 2025

- Implemented featurization approaches (STFT, Mel-spectrogram) within a Band-Split RNN to isolate vocal tracks
- Integrated HiFi-GAN for Mel-spectrogram inversion, improving audio reconstruction quality by reducing artifacts
- Optimized training via PyTorch AMP, hyperparameter tuning, and data augmentation on the MUSDB18 dataset

Explicit/Implicit Heap Allocator

Unix, C

May 2024 – Jun 2024

- Implemented the “malloc”, “realloc”, and “free” functions optimizing for request throughput and memory utilization
- Incorporated an explicit list of nodes to assign optimal locations for new memory requests and lower memory fragmentation
- Achieved 91% memory utilization via testing on heap activity memory requests from Emacs, Cmake, and Firefox

TECHNICAL SKILLS

Languages: Python, LaTeX, TypeScript, JavaScript, Lean, SQL, C++, C, HTML/CSS

Frameworks/Libraries: Pandas, NumPy, Matplotlib, scikit-learn, React, React Native, Next.js

Developer Tools: Git, Unix, Vim, VS Code, Apache Spark, Snowflake, Jupyter Notebook, Qt Creator